

CSE 344 Computer Vision Homework 15

Name: Arka Sarkar

Roll Number : 20182222

```
In [73]: import numpy as np
import tensorflow as tf
import pandas as pd
from keras import layers
from keras.layers import Input, Add, Dense, Activation, ZeroPadding2D, BatchNormalization, Flatten, Conv2D, AveragePooling2D, MaxPooling2D, GlobalMaxPooling2D
import matplotlib.pyplot as plt
from keras.models import Model, load_model
```

Reading the Dataset

```
In [2]: train_df = pd.read_csv("dataset/mnist_train.csv")
test_df = pd.read_csv("dataset/mnist_test.csv")

In [3]: train_df

Out[3]:
```

	label	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	...	28x19	28x20	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
0	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
...
59995	8	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59996	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59997	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59998	6	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59999	8	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

60000 rows x 785 columns

```
In [4]: train = np.array(train_df)
test = np.array(test_df)

Y_train = np.zeros((train.shape[0], 10))
X_train = []
Y_test = np.zeros((test.shape[0], 10))
X_test = []

for i in range(train.shape[0]):
    X_train.append(train[i,:].reshape((28,28,1)))
    Y_train[i,train[i,0]] = 1
for i in range(test.shape[0]):
    X_test.append(test[i,:].reshape((28,28,1)))
    Y_test[i,test[i,0]] = 1

X_train = np.array(X_train)
X_test = np.array(X_test)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)

(10000, 28, 28, 1)
(60000, 28, 28, 1)
(10000, 10)
(60000, 10)

In [8]: idx = 10
sample = X_train[idx]
print(Y_train[idx])
plt.imshow(sample, cmap = "gray")
[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]

Out[8]: <matplotlib.image.AxesImage at 0x7f2bd3655190>
```

```
In [19]: def LeNet5(input_shape = (28,28,1), stride = [1,1], filters = [5,5], pooling = [2,2], activation = ["relu", "relu", "relu", "relu"], n_filters = [6,16], n_layers = [120,84], pooling_technique = "max"):

    s1,s2 = stride
    f1,f2 = filters
    p1,p2 = pooling

    a1,a2,a3,a4 = activation
    n1, n2 = n_filters
    n1l, n12 = n_layers

    X_input = Input(input_shape)
    X = ZeroPadding2D((2,2))(X_input)

    X = Conv2D(filters = n1l, kernel_size = (f1,f1), strides = (s1,s1), activation = a1, input_shape = (32,32,1))(X)
    if(pooling_technique == "max"):
        X = MaxPooling2D(( p1, p1), strides=(2, 2))(X)
    elif(pooling_technique == "avg"):
        X = AveragePooling2D(( p1, p1), strides=(2, 2))(X)
    else:
        raise Exception("Wrong Pooling technique")

    X = Conv2D(filters = n2, kernel_size = (f2,f2), strides = (s2,s2), activation = a2)(X)
    if(pooling_technique == "max"):
        X = MaxPooling2D(( p2, p2), strides=(2, 2))(X)
    elif(pooling_technique == "avg"):
        X = AveragePooling2D(( p2, p2), strides=(2, 2))(X)
    else:
        raise Exception("Wrong Pooling technique")

    X = Flatten()(X)
    X = Dense(n1l, activation = a3)(X)
    X = Dense(n12, activation = a4)(X)
    X = Dense(10, activation = 'softmax')(X)

    model = Model(inputs = X_input, outputs = X, name='LeNet5')

    return model
```

Original LeNet 5

```
In [20]: model = LeNet5()
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train, Y_train, epochs = 10, batch_size = 512)
preds = model.evaluate(X_test, Y_test)
preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

Model: "LeNet5"
```

Layer (type)	Output Shape	Param #
input_7 (InputLayer)	[(None, 28, 28, 1)]	0
zero_padding2d_6 (ZeroPaddin	(None, 32, 32, 1)	0
conv2d_12 (Conv2D)	(None, 28, 28, 6)	156
max_pooling2d_12 (MaxPooling	(None, 14, 14, 6)	0
conv2d_13 (Conv2D)	(None, 10, 10, 16)	2416
max_pooling2d_13 (MaxPooling	(None, 5, 5, 16)	0
flatten_6 (Flatten)	(None, 400)	0
dense_18 (Dense)	(None, 120)	48120
dense_19 (Dense)	(None, 84)	10164
dense_20 (Dense)	(None, 10)	850
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		
Epoch 1/10		
118/118 [=====]	- 1s 3ms/step - loss: 5.1802 - accuracy: 0.6288	
Epoch 2/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.2296 - accuracy: 0.9339	
Epoch 3/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1342 - accuracy: 0.9612	
Epoch 4/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0882 - accuracy: 0.9730	
Epoch 5/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0711 - accuracy: 0.9779	
Epoch 6/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0551 - accuracy: 0.9827	
Epoch 7/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0430 - accuracy: 0.9866	
Epoch 8/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0351 - accuracy: 0.9888	
Epoch 9/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0252 - accuracy: 0.9924	
Epoch 10/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0205 - accuracy: 0.9941	
313/313 [=====]	- 1s 1ms/step - loss: 0.0733 - accuracy: 0.9809	
313/313 [=====]	- 0s 1ms/step - loss: 0.0733 - accuracy: 0.9809	
Loss = 0.07327043265104291		
Test Accuracy = 0.98089998960495		

Modified LeNet5

- filter size
- Stride
- Pooling technique. For example, you can try pooling by computing the "determinant" value.
- Activation function. For example, you can try out some new activation functions like "x.sigmoid(x)".
- Number of layers
- Number of filters

Changing Filter Size

```
In [21]: model = LeNet5(filters = [5,7])
model.summary()
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train, Y_train, epochs = 10, batch_size = 512)
preds = model.evaluate(X_test, Y_test)
preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

Model: "LeNet5"
```

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 28, 28, 1)]	0
zero_padding2d_7 (ZeroPaddin	(None, 32, 32, 1)	0
conv2d_14 (Conv2D)	(None, 28, 28, 6)	156
max_pooling2d_14 (MaxPooling	(None, 14, 14, 6)	0
conv2d_15 (Conv2D)	(None, 8, 8, 16)	4720
max_pooling2d_15 (MaxPooling	(None, 4, 4, 16)	0
flatten_7 (Flatten)	(None, 256)	0
dense_21 (Dense)	(None, 120)	30840
dense_22 (Dense)	(None, 84)	10164
dense_23 (Dense)	(None, 10)	850
Total params: 46,730		
Trainable params: 46,730		
Non-trainable params: 0		
Epoch 1/10		
118/118 [=====]	- 1s 3ms/step - loss: 4.7279 - accuracy: 0.5962	
Epoch 2/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.2628 - accuracy: 0.9255	
Epoch 3/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1423 - accuracy: 0.9572	
Epoch 4/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1082 - accuracy: 0.9683	
Epoch 5/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0839 - accuracy: 0.9747	
Epoch 6/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0579 - accuracy: 0.9816	
Epoch 7/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0534 - accuracy: 0.9826	
Epoch 8/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0406 - accuracy: 0.9873	
Epoch 9/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0314 - accuracy: 0.9905	
Epoch 10/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0266 - accuracy: 0.9917	
313/313 [=====]	- 1s 2ms/step - loss: 0.0904 - accuracy: 0.9736	
313/313 [=====]	- 1s 2ms/step - loss: 0.0904 - accuracy: 0.9736	
Loss = 0.09038401395082474		
Test Accuracy = 0.9735999703407288		

Changing Stride Size

```
In [22]: model = LeNet5(stride = [2,2])
model.summary()
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train, Y_train, epochs = 10, batch_size = 512)
preds = model.evaluate(X_test, Y_test)
preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

Model: "LeNet5"
```

Layer (type)	Output Shape	Param #
input_9 (InputLayer)	[(None, 28, 28, 1)]	0
zero_padding2d_8 (ZeroPaddin	(None, 32, 32, 1)	0
conv2d_16 (Conv2D)	(None, 14, 14, 6)	156
max_pooling2d_16 (MaxPooling	(None, 7, 7, 6)	0
conv2d_17 (Conv2D)	(None, 2, 2, 16)	2416
max_pooling2d_17 (MaxPooling	(None, 1, 1, 16)	0
flatten_8 (Flatten)	(None, 16)	0
dense_24 (Dense)	(None, 120)	2040
dense_25 (Dense)	(None, 84)	10164
dense_26 (Dense)	(None, 10)	850
Total params: 15,626		
Trainable params: 15,626		
Non-trainable params: 0		
Epoch 1/10		
118/118 [=====]	- 1s 3ms/step - loss: 4.3474 - accuracy: 0.2799	
Epoch 2/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.6701 - accuracy: 0.7772	
Epoch 3/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.3617 - accuracy: 0.8851	
Epoch 4/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.2577 - accuracy: 0.9204	
Epoch 5/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.2088 - accuracy: 0.9354	
Epoch 6/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1796 - accuracy: 0.9437	
Epoch 7/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1694 - accuracy: 0.9470	
Epoch 8/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1489 - accuracy: 0.9532	
Epoch 9/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1304 - accuracy: 0.9594	
Epoch 10/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1199 - accuracy: 0.9623	
313/313 [=====]	- 1s 2ms/step - loss: 0.1240 - accuracy: 0.9621	
313/313 [=====]	- 0s 1ms/step - loss: 0.1240 - accuracy: 0.9621	
Loss = 0.124035645103776		
Test Accuracy = 0.9621000289916992		

Changing Pooling Technique

```
In [23]: model = LeNet5(pooling_technique="avg")
model.summary()
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train, Y_train, epochs = 10, batch_size = 512)
preds = model.evaluate(X_test, Y_test)
preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

Model: "LeNet5"
```

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 28, 28, 1)]	0
zero_padding2d_9 (ZeroPaddin	(None, 32, 32, 1)	0
conv2d_18 (Conv2D)	(None, 28, 28, 6)	156
average_pooling2d (AveragePo	(None, 14, 14, 6)	0
conv2d_19 (Conv2D)	(None, 10, 10, 16)	2416
average_pooling2d_1 (Average	(None, 5, 5, 16)	0
flatten_9 (Flatten)	(None, 400)	0
dense_27 (Dense)	(None, 120)	48120
dense_28 (Dense)	(None, 84)	10164
dense_29 (Dense)	(None, 10)	850
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		
Epoch 1/10		
118/118 [=====]	- 1s 4ms/step - loss: 2.5328 - accuracy: 0.6309	
Epoch 2/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1370 - accuracy: 0.9586	
Epoch 3/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.0798 - accuracy: 0.9756	
Epoch 4/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0604 - accuracy: 0.9821	
Epoch 5/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0459 - accuracy: 0.9854	
Epoch 6/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0388 - accuracy: 0.9884	
Epoch 7/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0304 - accuracy: 0.9905	
Epoch 8/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0273 - accuracy: 0.9918	
Epoch 9/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0232 - accuracy: 0.9932	
Epoch 10/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0197 - accuracy: 0.9938	
313/313 [=====]	- 1s 2ms/step - loss: 0.0462 - accuracy: 0.9847	
313/313 [=====]	- 1s 2ms/step - loss: 0.0462 - accuracy: 0.9847	
Loss = 0.046213697642087936		
Test Accuracy = 0.9847000241279602		

Changing Activation Function

```
In [24]: model = LeNet5(activation = ["sigmoid", "sigmoid", "sigmoid", "sigmoid"])
model.summary()
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train, Y_train, epochs = 10, batch_size = 512)
preds = model.evaluate(X_test, Y_test)
preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

Model: "LeNet5"
```

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 28, 28, 1)]	0
zero_padding2d_10 (ZeroPaddi	(None, 32, 32, 1)	0
conv2d_20 (Conv2D)	(None, 28, 28, 6)	156
max_pooling2d_18 (MaxPooling	(None, 14, 14, 6)	0
conv2d_21 (Conv2D)	(None, 10, 10, 16)	2416
max_pooling2d_19 (MaxPooling	(None, 5, 5, 16)	0
flatten_10 (Flatten)	(None, 400)	0
dense_30 (Dense)	(None, 120)	48120
dense_31 (Dense)	(None, 84)	10164
dense_32 (Dense)	(None, 10)	850
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		
Epoch 1/10		
118/118 [=====]	- 1s 4ms/step - loss: 2.2663 - accuracy: 0.1697	
Epoch 2/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.8766 - accuracy: 0.7940	
Epoch 3/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.3179 - accuracy: 0.9279	
Epoch 4/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1909 - accuracy: 0.9515	
Epoch 5/10		
118/118 [=====]	- 0s 4ms/step - loss: 0.1383 - accuracy: 0.9626	
Epoch 6/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.1168 - accuracy: 0.9673	
Epoch 7/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0987 - accuracy: 0.9719	
Epoch 8/10		
118/118 [=====]	- 0s 3ms/step - loss: 0.0866 - accuracy: 0.9750	
Epoch 9/10		
118/118 [